

# Sai Ravi Teja Gangavarapu

(412) 251-7161 | [sairavig@andrew.cmu.edu](mailto:sairavig@andrew.cmu.edu) | Pittsburgh, PA | [floaredor.vercel.app](https://floaredor.vercel.app)

## EDUCATION

### Carnegie Mellon University

Master of Software Engineering - Scalable Systems (Statistics, Prod Mgmt, Agile Methods, **ML**)

Pittsburgh, PA

Aug 2025 – Dec 2026

### University of Florida

Senior Cert + Master's coursework in Computer Science (Algorithms, UX, Graphics - **4.0**)

Gainesville, FL

Jan 2024 – May 2024

## SKILLS

Python, MATLAB, C++, Go, JavaScript, TypeScript, Java, Rust, CUDA, PyTorch, TensorFlow, FastAPI, REST API, React, React Native, Next.js, Langchain, Langflow, Docker, Kubernetes, AWS, Git, PostgreSQL, MySQL, NoSQL, MongoDB, OpenGL, Spark, Pandas, Numpy, Pydantic, SQLAlchemy, React Query, Tanstack, Zustand, WebSockets, Redis, Celery, ArgoCD, SQL, HTML, Agile, JIRA, (Deep Learning, Machine Learning, Probability and Statistics, Linear Algebra, DSA, MathAcademy ML Mathematics)

## EXPERIENCE

### Founding Software Engineer — Tapsta

Jul 2024 – Jun 2025

- Architected end-to-end full-stack social rewards application using React, React Native, FastAPI, and PostgreSQL.
- Designed database architecture with 25+ schemas and optimized client-side operations with caching and debouncing, achieving 30% faster response times, serving over 1500 active users.
- Orchestrated Plaid API integration alongside REST API creation, establishing secure bank connectivity and enabling automated cashback processing, plus ACH transfers for over 1,500 application users.
- Made containerized CI/CD pipeline with Docker, AWS, and ArgoCD that reduced deployment time by 70%.
- Led development team of 3 SDE interns, establishing code review processes that contributed to **30%** faster feature delivery.
- Built a semantic people search engine for alumni networks using embedding-based similarity using LangGraph.

### Research Assistant - Fan Lab @UF

Feb 2024 – Jun 2024

- Applied genomic foundational models for rare disease prediction using **transformer architectures**, processing **15,000+ DNA sequences**. Developed multi-layer perception models using protein embeddings (**ESM3**) for rare disease prediction.

### SDE Intern — Catalog.fi

Apr 2023 – Dec 2023

- Implemented full-stack crypto analytics platform using React, FastAPI, Postgres, Pandas, and Golang with real-time monitoring dashboard and a microservice for a leaderboard system for garden.finance, boosting user engagement by 30%.
- Improved backend performance using Go and PostgreSQL optimizations, supporting \$150M+ trading volume over 30 days with 20% faster queries on garden.finance.
- Taught LLMs to use APIs, orchestrating fine-tuning of LLaMA and BERT models using PyTorch, achieving 90% accuracy in translating natural language to API calls and reducing API integration time by 25%.

### Research Assistant - Mahindra Ecole Centrale

Dec 2022 – Jan 2025

- Published research paper on emotion-targeted music generation using FFT, evolutionary algorithms and self-organizing maps (IEEE CEC 24). Implemented ALL-GAN model with t-SNE and PCA for clustering analysis on 1000+ samples, achieving unsupervised music genre classification. Also, built MIR pipelines extracting 100+ time- and frequency-domain features.

## PROJECTS

### Project RECON | Raspberry Pi, OpenMPI, GlusterFS | \$2000 funding

- Architected 8-node Raspberry Pi 4B **compute cluster** with distributed storage, LDAP authentication, and Slurm job scheduling, serving 400+ students for coursework.

### OneAIClick.com | Python, React.js, HuggingFace, PEFT/LoRA (Built no-code/low code tool)

- Developed an LLM fine-tuning abstraction tool for HuggingFace models that cut development cycles by 60% and reduced boilerplate code by 75%, enabling rapid idea validation and full data privacy.

### CUDA Ray Tracer & Audio-Reactive Visualizer | C++, CUDA, OpenGL

- Built GPU-accelerated ray tracing system achieving **1,600x speedup** (7.5s to 0.0045s per frame) over CPU implementation, with real-time audio analysis and 256-agent boids flocking synchronized to music frequency bands.

### RainPod: AI-Powered Interactive Podcast | Agents, CrewAI, Python, FastAPI, React, Nextjs

- Engineered multi-agent system converting PDFs/topics into interactive AI podcasts for personalised educational content.

## LEADERSHIP AND ACHIEVEMENTS

---

**President, Enigma, the Computer Science Club (2021-2023):** Orchestrated Generative AI Hackathon with **NVIDIA** and Game Jam with Ubisoft, conducted technical workshops on Gamedev, ML, Linux, heading outreach initiatives increasing the club membership by 40%, reaching over 2000 students.

**1st Place WaffleHacks 2024** (320+ participants): Built FlashFocus - a Chrome extension to block distractions and turn them into learning with AI flashcard quizzes using RAG system. (React, FastAPI, SQL, Groq)

**1st Place Talentmapp Hack4Hire 2023** (300+ participants): Made an intelligent task-tracking application with LLM-powered suggestions (MongoDB, Flask, React.js, ChatGPT)

**1st Place Noderunner Hackathon 2023** (50 teams): Engineered distributed cluster implementing Raft consensus protocol from scratch (Python, Flask, Multithreading)

**Publications: IEEE CEC 2024** - "Emotion Aligned Music Composition from Sound Fundamentals Using Differential Evolution"